

## Faster R-CNN + Inception v2

### Coco metrics:

mAP@.50IOU: 0.498137

classification\_loss: 0.034073

localization\_loss: 0.034354

objectness\_loss: 0.079878

### PASCAL VOC metrics:

AP: 68.10% (Green)





AP: 60.98% (Red)

AP: 65.43% (Yellow)

AP: 0.05% (off)

mAP: 48.64%



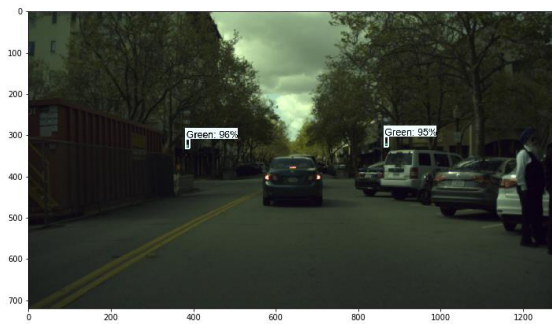
### Precision x Recall curves at 0.5 IoU:



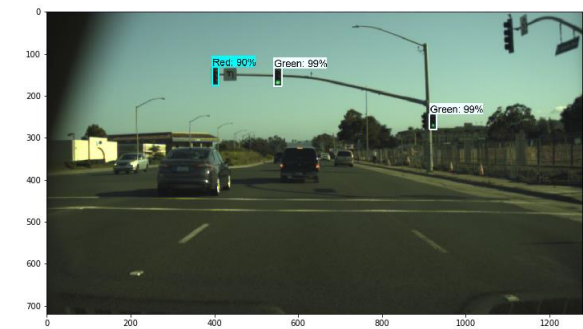
Trained the model for 100k iterations

## Detections:

### Test set images



### Train set images





Tested on a random video feed



## Faster R-CNN + Resnet101

### PASCAL VOC metrics:

AP: 61.30% (Green)





AP: 61.22% (Red)

AP: 57.75% (Yellow)

AP: 0.01% (off)

mAP: 45.07%

### Precision x Recall curves at 0.5 IoU:



Trained the model for 50k iterations



## Detections:

### Test set images



### Train set images



Random video





